

```
In [73]: import pandas as pd
import seaborn as sns
```

```
In [74]: df = pd.read_csv("D:\c drive settinf\Desktop\datasets\Churn_Modelling.csv")
df.head()
```

```
Out[74]:
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age | Tenure | Balance |
|---|-----------|------------|----------|-------------|-----------|--------|-----|--------|-----------|
| 0 | 1 | 15634602 | Hargrave | 619 | France | Female | 42 | 2 | 0.00 |
| 1 | 2 | 15647311 | Hill | 608 | Spain | Female | 41 | 1 | 83807.86 |
| 2 | 3 | 15619304 | Onio | 502 | France | Female | 42 | 8 | 159660.80 |
| 3 | 4 | 15701354 | Boni | 699 | France | Female | 39 | 1 | 0.00 |
| 4 | 5 | 15737888 | Mitchell | 850 | Spain | Female | 43 | 2 | 125510.82 |

```
In [75]: df.shape
```

```
Out[75]: (10000, 14)
```

```
In [76]: df.describe()
```

```
Out[76]:
```

| | RowNumber | CustomerId | CreditScore | Age | Tenure | Balance | Num |
|-------|-------------|--------------|--------------|--------------|--------------|---------------|--------------|
| count | 10000.00000 | 1.000000e+04 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 |
| mean | 5000.50000 | 1.569094e+07 | 650.528800 | 38.921800 | 5.012800 | 76485.889288 | |
| std | 2886.89568 | 7.193619e+04 | 96.653299 | 10.487806 | 2.892174 | 62397.405202 | |
| min | 1.00000 | 1.556570e+07 | 350.000000 | 18.000000 | 0.000000 | 0.000000 | |
| 25% | 2500.75000 | 1.562853e+07 | 584.000000 | 32.000000 | 3.000000 | 0.000000 | |
| 50% | 5000.50000 | 1.569074e+07 | 652.000000 | 37.000000 | 5.000000 | 97198.540000 | |
| 75% | 7500.25000 | 1.575323e+07 | 718.000000 | 44.000000 | 7.000000 | 127644.240000 | |
| max | 10000.00000 | 1.581569e+07 | 850.000000 | 92.000000 | 10.000000 | 250898.090000 | |

```
In [77]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname                10000 non-null  object
 3   CreditScore            10000 non-null  int64
 4   Geography              10000 non-null  object
 5   Gender                 10000 non-null  object
 6   Age                   10000 non-null  int64
 7   Tenure                 10000 non-null  int64
 8   Balance                10000 non-null  float64
 9   NumOfProducts          10000 non-null  int64
10   HasCrCard              10000 non-null  int64
11   IsActiveMember         10000 non-null  int64
12   EstimatedSalary        10000 non-null  float64
13   Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [78]: df.isnull().sum()
```

```
Out[78]: RowNumber      0
CustomerId    0
Surname       0
CreditScore   0
Geography     0
Gender        0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
```

```
In [79]: df.duplicated().sum()
```

```
Out[79]: 0
```

```
In [80]: df["Geography"].value_counts()
```

```
Out[80]: France      5014
Germany    2509
Spain      2477
Name: Geography, dtype: int64
```

```
In [81]: df["Gender"].value_counts()
```

```
Out[81]: Male        5457
Female      4543
Name: Gender, dtype: int64
```

```
In [82]: df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
```

```
In [83]: df = pd.DataFrame(df)
df
```

Out[83]:

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsAc |
|------|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|------|
| 0 | 619 | France | Female | 42 | 2 | 0.00 | 1 | 1 | |
| 1 | 608 | Spain | Female | 41 | 1 | 83807.86 | 1 | 0 | |
| 2 | 502 | France | Female | 42 | 8 | 159660.80 | 3 | 1 | |
| 3 | 699 | France | Female | 39 | 1 | 0.00 | 2 | 0 | |
| 4 | 850 | Spain | Female | 43 | 2 | 125510.82 | 1 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 771 | France | Male | 39 | 5 | 0.00 | 2 | 1 | |
| 9996 | 516 | France | Male | 35 | 10 | 57369.61 | 1 | 1 | |
| 9997 | 709 | France | Female | 36 | 7 | 0.00 | 1 | 0 | |
| 9998 | 772 | Germany | Male | 42 | 3 | 75075.31 | 2 | 1 | |
| 9999 | 792 | France | Female | 28 | 4 | 130142.79 | 1 | 1 | |

10000 rows × 11 columns

In [84]: `# extract the data from the dataset`
`categorical_columns = df.select_dtypes(include=["object"]).columns.tolist()`

categorical_columns

In [85]: categorical_columns

Out[85]: ['Geography', 'Gender']

In [86]: `# import sklearn and onehotencoder`
`from sklearn.preprocessing import OneHotEncoder`

In [87]: `# initialize the onehotencoding`
`ohe = OneHotEncoder(sparse=False)`

In [88]: `one_hot_encoded = ohe.fit_transform(df[categorical_columns])`

In [89]: `# creating dataframe for onehotencoded.`
`one_hot_df = pd.DataFrame(one_hot_encoded)`

In [90]: `# to get the column name of the we use the get_feature_names_out()`
`columns_name = ohe.get_feature_names_out(categorical_columns)`

In [91]: `# concatenate the ohe_dataframe with original one`
`df_encoded = pd.concat([df, one_hot_df], axis=1)`

In [92]: `df_encoded_1 = df_encoded.drop(categorical_columns, axis=1)`

In [96]: `df_encoded_1.columns`

```
Out[96]: Index([    'CreditScore',          'Age',          'Tenure',
          'Balance',    'NumOfProducts',    'HasCrCard',
          'IsActiveMember', 'EstimatedSalary',    'Exited',
          0,          1,          2,
          3,          4],
          dtype='object')
```

```
In [98]: # import libraries for performing the train_test_split
from sklearn.model_selection import train_test_split
```

```
In [101... x_train,x_test,y_train,y_test = train_test_split(df_encoded_1.drop(columns=['Exited',
```

```
In [102... x_train
```

```
Out[102]:
```

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-------------|-----|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 4322 | 508 | 31 | 8 | 72541.48 | 1 | 1 | 0 | 162454.35 | 0 |
| 955 | 706 | 44 | 4 | 129605.99 | 1 | 0 | 0 | 60131.43 | 0 |
| 4019 | 620 | 31 | 2 | 166833.86 | 2 | 1 | 1 | 137616.78 | 0 |
| 8314 | 643 | 33 | 4 | 0.00 | 2 | 1 | 1 | 151604.92 | 0 |
| 9272 | 739 | 42 | 2 | 141642.92 | 2 | 1 | 0 | 177650.00 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9841 | 567 | 46 | 1 | 68238.51 | 2 | 1 | 1 | 109929.00 | 0 |
| 3352 | 591 | 40 | 2 | 99886.42 | 2 | 1 | 1 | 83001.57 | 0 |
| 1645 | 506 | 41 | 3 | 57745.76 | 1 | 1 | 0 | 113612.78 | 0 |
| 7531 | 692 | 40 | 6 | 163505.16 | 1 | 0 | 0 | 92319.67 | 0 |
| 7554 | 706 | 30 | 6 | 87609.68 | 2 | 0 | 0 | 138615.00 | 0 |

8000 rows × 13 columns

```
In [103... x_test
```

Out[103]:

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|------|-------------|-----|--------|-----------|---------------|-----------|----------------|-----------------|
| 6759 | 705 | 92 | 1 | 126076.24 | 2 | 1 | 1 | 100000 |
| 2177 | 667 | 24 | 4 | 0.00 | 2 | 0 | 1 | 100000 |
| 818 | 497 | 27 | 9 | 75263.16 | 1 | 1 | 1 | 100000 |
| 5655 | 695 | 63 | 1 | 146202.93 | 1 | 1 | 1 | 100000 |
| 8531 | 723 | 30 | 1 | 0.00 | 3 | 1 | 0 | 100000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 886 | 739 | 38 | 0 | 128366.44 | 1 | 1 | 0 | 100000 |
| 2900 | 626 | 26 | 8 | 148610.41 | 3 | 0 | 1 | 100000 |
| 2717 | 775 | 70 | 6 | 119684.88 | 2 | 1 | 1 | 100000 |
| 7672 | 555 | 30 | 1 | 0.00 | 2 | 0 | 0 | 100000 |
| 7838 | 569 | 32 | 8 | 145330.43 | 1 | 1 | 1 | 100000 |

2000 rows × 13 columns

```

In [104... # feature scaling
from sklearn.tree import DecisionTreeClassifier

In [105... clf = DecisionTreeClassifier(max_depth=5, random_state=42)

In [108... clf.fit(x_train,y_train)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.
  warnings.warn(
Out[108]: DecisionTreeClassifier(max_depth=5, random_state=42)

In [109... predict = clf.predict(x_test)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.
  warnings.warn(

In [111... from sklearn.metrics import accuracy_score, classification_report
print("Accuracy:", accuracy_score(y_test,predict)*100)

Accuracy: 84.39999999999999

In [ ]:

```